



Relative BPB efficiency and CORE performance by model scale

● Pythia (GPT-NeoX, 2023) ● SmoLM2 (Llama, 2024) ○ Open = best checkpoint shown (see caption)



Both panels: ↑ = candidate beats its size-matched GPT-2 (identical pipeline). Pythia sits at/below parity on both metrics; SmoLM2 stays at BPB parity or below (left) while showing a small, consistent downstream edge on CORE (right). CORE (limit=500) is secondary to BPB and noisy; ± 1 stderr shown.