


Exp B — a metric disagreement: BPB efficiency is not the same as downstream capability



Both panels: ↑ = candidate beats its size-matched GPT-2 (identical pipeline). Pythia sits at/below parity on BOTH metrics. SmolLM2 never beats GPT-2 on BPB (left) yet shows a small, consistent downstream edge on CORE (right) — the two measures disagree. CORE (limit=500) is secondary and noisy; ± 1 stderr shown.