


# Corpus compute-equivalent gain vs. OWT reference



124M GPT-2 baseline scale. Each corpus's total compute-equivalent gain vs. the old-recipe-OWT reference arm (1x, by definition), under each recipe separately. Marker shape = recipe (circle = old GPT-2, square = current training). Hollow markers below the 1x line did not reach the threshold within their compute budget — a bound, not a measured value.