


Corpus-only compute-equivalent gain, training recipe held fixed



124M GPT-2 baseline scale. Each point swaps ONLY the training corpus (OWT → the labeled corpus) while holding the training recipe/algorithm fixed at the value given by its marker shape — this isolates the corpus's own contribution, with no algorithm effect mixed in. RefinedWeb/DCLM give a different corpus multiplier under each recipe (old GPT-2 recipe ~3.3–3.5x vs. current training recipe ~1.6x): the corpus and training-recipe interventions interact, so there is no single recipe-independent "value of the corpus".