




Exp A — corpus intervention at the 124M GPT-2 baseline scale (neutral eval: wiki_eval_union)

● old-algo (GPT-2) recipe ■ current-arch (modded) recipe ○ hollow = censored (bound, not a value)



Panel C isolates the corpus lever within each fixed recipe. The two recipes give different data multipliers (old-algo ~3.3–3.5x vs current-arch ~1.6x on RefinedWeb/DCLM): the corpus and training-recipe interventions INTERACT, so a Shapley attribution of “how much the data is worth” is context-dependent and not a single number. C4 (2020) never reaches the threshold under either recipe — corpus quality is non-monotonic in release date.