


Corpus compute-equivalent gain at the 124M GPT-2 baseline scale



Held-out, decontaminated Wikipedia evaluation set ("wiki_eval_union"). Panel C holds the recipe fixed and swaps only the training corpus (OWT → corpus); the two recipes give different corpus multipliers on RefinedWeb/DCLM (old GPT-2 recipe ~3.3–3.5x vs current training recipe ~1.6x) — see caption for interpretation.