


Neutral-corpus BPB vs. GPU-hours by training corpus



Same architecture, four training corpora, vs its fixed size-matched GPT-2-OWT bar. RefinedWeb (2023) and DCLM (2024) cross under both architectures; OWT (2019) and C4 (2020) stay censored under both. “Censored” = the arm never reaches the threshold within its compute budget, so its compute-to-threshold is a bound, not a value (hollow markers).