


Exp B — a large speedrun-style algorithmic advantage does NOT reproduce across six matched comparisons



Each modern architecture vs a GPT-2 trained through the identical pipeline at the same size (OWT, 8.87B tokens, converged 512k-batch recipe). Lower BPB is better; no candidate reaches its matched GPT-2 bar. “Censored” = the arm never reaches the threshold within its compute budget, so its compute-to-threshold is a bound, not a value (hollow markers).