


Compute to reach the neutral-BPB threshold



The primitive behind every compute-equivalent gain: train to a fixed neutral-BPB threshold and read off the GPU-hours it took (fewer = more compute-efficient). Worked example at the 124M GPT-2 baseline scale, current training recipe, old-algorithm/old-data arm.