


Compute to reach the neutral-BPB threshold

124M GPT-2 baseline scale · current training recipe



The primitive behind every compute-equivalent gain: train to a fixed neutral-BPB threshold and read off the GPU-hours it took (fewer = more compute-efficient).