


Compute-equivalent gain vs. GPT-2 baseline scale



X-axis: dimensions of the reference GPT-2 model, not the trained candidate's own parameter count.