


Compute-equivalent-gain multipliers vs model scale



x-axis = GPT-2 baseline scale (baseline dimensions), not candidate parameter count: current-arch carries ~498.8M params at the 124M-baseline dimensions (the value-embed / U-net additions are the algorithm being measured). Eval set: wiki_eval. Lineages use independent hardware GPU-hour bases and are never mixed in one ratio.